

UNIVERSITY TUNKU ABDUL RAHMAN

LEE KONG CHIAN FACULTY OF ENGINEERING AND SCIENCE

UECS3483/ UECS3213/ UECS3453 DATA MINING

JANUARY 2024 TRIMESTER

LECTURER: SU LEE SENG

No.	Student Name	Student ID	Programme	Practical Group
1.	Kim Bei Er	2103893	SE	P1
2.	Kuik Rui Luan	2103790	SE	P1
3.	Lim Juan Hong	2105435	SE	P1


Table of Contents

1. Data Exploration	4
1.1 Variable Summarization	4
1.2 Objective, Potential Issue and Limitation	5
2. Data Description	6
2.1 Descriptive Statistics	6
3. Data Preprocessing	12
3.1 Data Cleaning	12
3.2 Feature Selection	12
3.3 Data Splitting	13
3.4 Detect and Remove Outlier	13
3.5 Data Transformation and Standardization	13
4. Data Mining Model	15
4.1 Model Training (Without Resampling)	15
4.2 Data Resampling Technique	17
4.3 Model after Resampling and Hyperparameter Tuning	18
4.3.1 Logistic Regression	18
4.3.2 Naïve Bayes	19
4.3.3 Decision Tree	20
4.3.4 Random Forest	22
4.3.5 K-Nearest Neighbors (KNN)	24
4.3.6 Support Vector Machine	25
4.3.7 Neural Network	26
4.4 Comparison and Selection	28
5. Interaction stage	29
6. References	31

README

Method 1 (Using Colab)

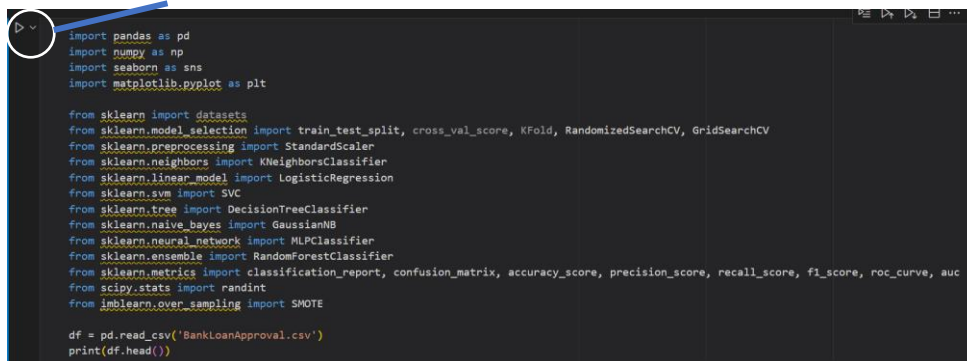
1. To run the code, kindly click the link below:
https://colab.research.google.com/drive/11tZIHh_TE18cNbZwuOjy7UDmwVv0deRe?usp=sharing
2. Upload 'BankLoanApproval.csv' and 'NewApplicants.csv' before running the code.




Click on the upload file icon to upload 'BankLoanApproval.csv' and 'NewApplicants.csv'

Method 2 (Using Visual Studio Code)


1. Open the 'Data Mining Coding.ipynb' in the submission folder using Visual Studio Code.
2. Click the 'Run' button



3. Select 'Python Environments'



4. Choose the python environment in your laptop.



1. Data Exploration

1.1 Variable Summarization

The “Bank Loan Approval” dataset provides information about individuals and their loans, with the objective to predict those at high risk of defaulting on their loans. This dataset comprises 18 key features that delve into various aspects of the applicant’s background and loan specifics:

1. **LoanID**: A unique identifier for each loan, used for tracking purposes.
2. **Age**: Age of the borrower at the time of loan application, indicating their level of financial responsibility and stability.
3. **Income**: Annual income of the borrower, help in assess their ability to repay the loan.
4. **LoanAmount**: Amount of money being borrowed.
5. **CreditScore**: Credit score of the borrower, indicating their creditworthiness.
6. **MonthsEmployed**: Number of months the borrower has been employed, indicating their stability in their job.
7. **NumCreditLines**: Number of credit lines the borrower has open, reflecting their credit behavior.
8. **InterestRate**: Interest rate charged on the loan amount.
9. **LoanTerm**: Term length of the loan in months.
10. **DTIRatio**: Debt-to-Income ratio, indicating the borrower’s debt compared to their income.
11. **Education**: Highest level of education attained by borrower (PhD, Master’s, Bachelor’s, High School), indicating their earning potential.
12. **EmploymentType**: Type of employment status of the borrower (Full-time, Part-time, Self-employed, Unemployed), which affect their ability to repay the loan.
13. **MaritalStatus**: Marital status of the borrower (Single, Married, Divorced), which reflect their financial burden and responsibility.
14. **HasMortgage**: Whether the borrower has a mortgage (Yes or No), which can affect their financial obligations.
15. **HasDependents**: Whether the borrower has dependents (Yes or No), which can affect their financial responsibilities.
16. **LoanPurpose**: Purpose of loan (Home, Auto, Education, Business, Other), indicating the borrower’s financial goals and intentions.

17. **HasCoSigner:** Whether the loan has a co-signer (Yes or No), which can affect the qualification for the loan.
18. **Default:** Binary target variable indicating whether the loan defaulted (1) or not (0), which is the variable to be predicted in the prediction model.

1.2 Objective, Potential Issue and Limitation

The objective of the “Bank Loan Approval” dataset is to develop a model to accurately classify the loan applications as ‘Default’ or ‘Non-default’ based on the application information. This prediction can help financial institutions or lenders make informed decisions about loan approvals to **minimize the financial losses** while **maximize the approval rate for deserving applicants**. One potential issue with this dataset in developing the machine learning model could be imbalanced classes, where the number of defaulters (1) is significantly lower than the number of non-defaulters. This imbalance could lead to biased models that are better at predicting non-defaulters than defaulters. However, there could be ethical considerations related to using certain features, such as age, education, or marital status, in predicting loan default risk, as they could lead to biased decisions.

2. Data Description


2.1 Descriptive Statistics

```
#Descriptive statistics
df.describe()
```

	Age	Income	LoanAmount	CreditScore	MonthsEmployed	NumCreditLines	InterestRate	LoanTerm	DTIRatio	Default
count	255327.000000	255327.000000	255327.000000	255327.000000	255327.000000	255327.000000	255327.000000	255327.000000	255327.000000	255327.000000
mean	43.498059	82500.225585	127579.236559	574.266125	59.542516	2.501036	13.492848	36.025896	0.500222	0.116118
std	14.990304	38963.150663	70841.308245	158.904496	34.643129	1.117021	6.636456	16.969297	0.230917	0.320367
min	18.000000	15000.000000	5000.000000	300.000000	0.000000	1.000000	2.000000	12.000000	0.100000	0.000000
25%	31.000000	48826.000000	66156.000000	437.000000	30.000000	2.000000	7.770000	24.000000	0.300000	0.000000
50%	43.000000	82467.000000	127557.000000	574.000000	60.000000	2.000000	13.460000	36.000000	0.500000	0.000000
75%	56.000000	116219.000000	188986.500000	712.000000	90.000000	3.000000	19.250000	48.000000	0.700000	0.000000
max	69.000000	149999.000000	249999.000000	849.000000	119.000000	4.000000	25.000000	60.000000	0.900000	1.000000

Descriptive Statistics

The descriptive statistics reveal valuable insights into the characteristics of the loan applicant dataset. On average, applicants are approximately 43.5 years old, with ages ranging from 18 to 69 years. The average income stands at \$82,500, indicating a moderate-income level. However, the income distribution may exhibit a slight right skew, as the mean income slightly surpasses the median. Applicants request loans with varying amounts, as reflected by the average loan amount of \$127,579 and a considerable standard deviation of \$70,841. This wide variability suggests diverse financial needs among applicants. The creditworthiness of applicants, assessed by credit scores, with an average score of 573, indicating a good credit standing. Employment history among applicants varies, with an average tenure of approximately 59.5 months. Despite most applicants work for more than 30 months, there are also applicants who have never been employed and they might be university students or unemployed fresh graduate. The number of credit lines held by applicants ranges from 1 to 4, with the majority possessing 2 credit lines. Interest rates offered to applicants vary widely, ranging from 2% to 25%, with an average rate of approximately 13.49%. Loan terms span from 12 to 60 months, accommodating diverse preferences and financial circumstances. Applicants exhibit a diverse range of debt-to-income (DTI) ratios, with an average ratio of 0.5. The DTI ratio serves as a crucial indicator of an applicant's financial ability to manage debt obligations. Regarding loan defaults, approximately 11.6% of applicants defaulted on their loans. This proportion suggests potential imbalances in the dataset's target variable, which warrant careful consideration during model development and evaluation.

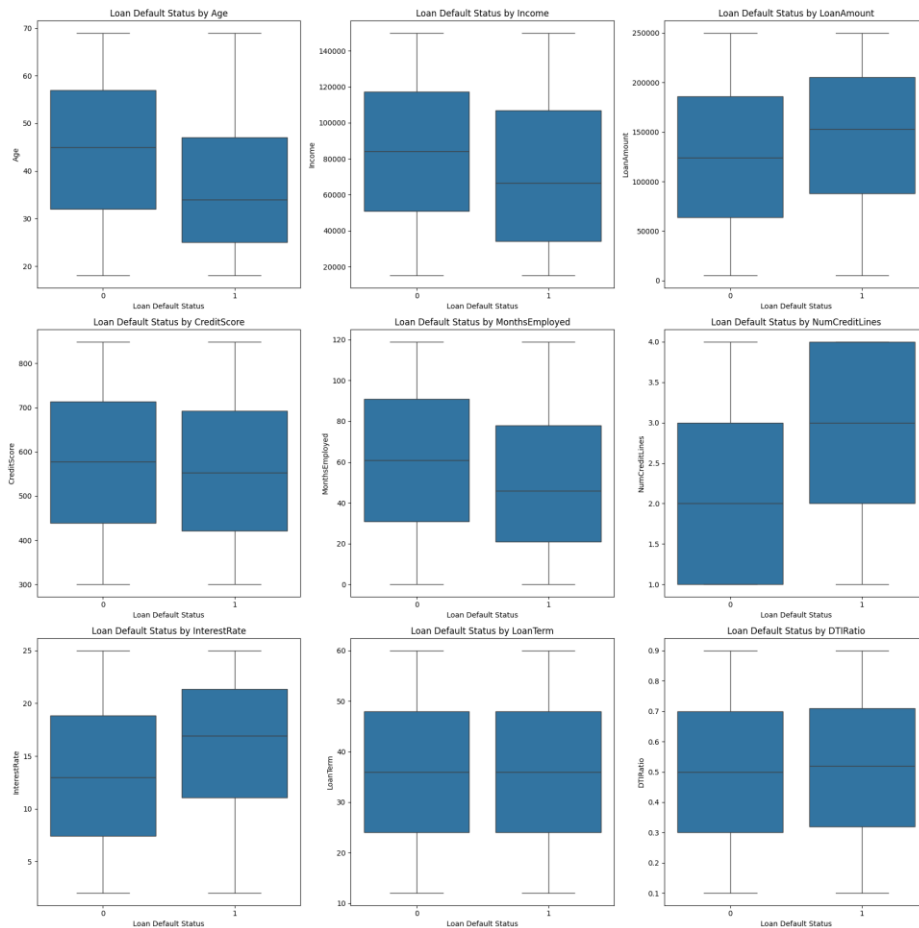


```
✓ 0s ▶ print("Number of class 0 (Non-default) instance: ", df[df["Default"] == 0].shape[0])
print("Number of class 1 (Default) instance: ", df[df["Default"] == 1].shape[0])

Number of class 0 (Non-default) instance: 225679
Number of class 1 (Default) instance: 29648
```

Count Plot for Default Variable


The count plot reveals that out of the total applicants, 225,679 individuals who did not default on their loans were approved for bank loans. Conversely, 29,648 applicants who defaulted on their loans had their bank loan requests rejected. This indicates the loan approval rate is relatively high. Moreover, the **ratio of class 0 ('non-default') to class 1 ('default')** is 7.6:1. The significant disparity in the number of instances between the two classes indicates a **class imbalance**. Class imbalance can potentially impact the performance of machine learning models, particularly those trained on imbalanced datasets, as they may become biased towards the majority class.



Boxplots for Numerical Variables

The boxplot reveals that **most younger applicants are more likely to default on their loans**. This may be because young applicants may not have ability to pay the loan. Applicants with **higher income get to repay for their bank loans** as this may be because of higher income individuals may have higher financial ability to pay back their loans. Applicants who **request for higher loan amount have higher probabilities on defaulting their loans** and this may pose higher risk to the bank. The credit score for both loan approval status is approximately the same but **non-defaulters have slightly higher credit scores** compared to defaulters, showcasing their higher creditworthiness. Applicants who have **worked for longer months are less likely to default on their loans**. This may be due to the applicants who have more working experience may have more income and more able to pay back the loan. **Non-defaulters mostly have less than 3 credit lines**, indicating they have less debt to pay. **Applicants who received lower interest rates** may have higher credit scores or better credit histories. Bank authorities may be more willing to approve loans for these applicants because they have demonstrated a lower likelihood of defaulting on their loans. The boxplots show similar distribution of loan terms for both ‘Default’ and ‘Non-default’ classes. This indicates


that the **length of the loan term may not influence the likelihood of default**. Thus, it may not be a significant factor in determining whether an applicant defaults on their loan. Applicants with **lower Debt-to-Income (DTI) Ratio tends to pay back for their loans** compared to applicants with higher DTI Ratio. This is because lower DTI Ratio indicates that the applicant has a higher capability to manage debt payment. In summary, the boxplots of loan default status by nine of the numerical variables show the distribution of data within loan default status of 0 ('non-default') and 1 ('default') and there is no outlier observed.



Bar Plots for Categorical Variables

The bar plots depicting loan default status across seven categorical variables reveal a consistent distribution of unique values within each category, regardless of the loan default outcome (0 for 'non-default' and 1 for 'default'). This uniform distribution indicates that the presence of **these categorical variables does not significantly influence the likelihood of loan default**. This observation suggests that the variables represented by these categories may not have a strong predictive power in determining loan default. However, it is important to further

investigate the relationship between these categorical variables and ‘Default’ variable through statistical tests or feature importance analysis to confirm their impact on the outcome.



Correlation Heatmap of Numerical Features

According to the correlation heatmap, there appears to be no significant correlation among the numerical features within the ‘Bank Loan Approval’ dataset, as indicated by null correlation coefficients. However, the Default variable shows some degree of correlation with other numerical variables. The **Default variable exhibits the strongest positive correlation (0.13) with the Interest Rate variable**, indicating that higher interest rates are associated with a higher likelihood of defaulting on loans. Additionally, there is a low positive correlation (0.09) between the Loan Amount variable and the Default variable, suggesting that larger loan amounts may receive rejection on bank loans due to the higher risk of default. Furthermore, the Default variable shows slight positive correlations with the Number of Credit Lines and DTI Ratio variables, implying that applicants with more credit lines or higher debt-to-income ratios may have a slightly increased possibility of defaulting.

Conversely, the **Age variable exhibits the strongest negative correlation (-0.17) with the Default variable**, indicating that older applicants may have a lower chance of defaulting on loans due to their financial stability. Similarly, the Income and Months Employed variables

demonstrate negative correlations (-0.10) with the Default variable, suggesting that higher incomes and longer employment durations may be associated with a reduced risk of default. The Credit Score variable shows a weak negative correlation (-0.03) with the Default variable, implying that higher credit scores may be associated with a lower likelihood of default. Interestingly, the Loan Term variable does not appear to have a significant correlation with the Default variable, suggesting that the length of the **loan term may not strongly influence the loan default status.**

3. Data Preprocessing

3.1 Data Cleaning

```
[6] #Data preprocessing for training set
#Check missing values for training set data
print(x_train.isnull().sum())

Age          0
Income       0
LoanAmount   0
CreditScore  0
MonthsEmployed 0
NumCreditLines 0
InterestRate 0
LoanTerm     0
DTIRatio     0
Education    0
EmploymentType 0
MaritalStatus 0
HasMortgage  0
HasDependents 0
LoanPurpose  0
HasCoSigner  0
dtype: int64
```

Check Missing Value

```
#Check duplicate rows for training set data
x_train.duplicated().any()

False
```

Check Duplicate Row

In this stage, we will handle the missing value and duplicate observation. In this case, the dataset does not have any missing value or duplicate observation. Thus, no action is needed for handling the missing value or duplicate row.

3.2 Feature Selection

```
Number of unique value in a variable
'LoanID':255327
'Age':52
'Income':114619
'LoanAmount':158725
'CreditScore':550
'MonthsEmployed':120
'NumCreditLines':4
'InterestRate':2301
'LoanTerm':5
'DTIRatio':81
'Education':4
'EmploymentType':4
'MaritalStatus':3
'HasMortgage':2
'HasDependents':2
'LoanPurpose':5
'HasCoSigner':2
'Default':2
```

Unique value in each variable

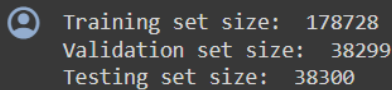
The 'LoanID' from the dataset appeared to be the unique value and it is useless in the analysis of data. In order to avoid the consumption of resource by irrelevant information that could impact the performance and accuracy, thus the **LoanID is dropped from dataset**.

```
#Define feature vector and target variable
x = df.drop(['LoanID', 'Default'], axis=1)
y = df['Default']
```

Define 'x' and 'y' Variables

'Default' variable is assigned to y as the target variable, while other features in DataFrame 'df' (except 'LoanID' and 'Default' variables) are assigned to x as independent variables.

3.3 Data Splitting



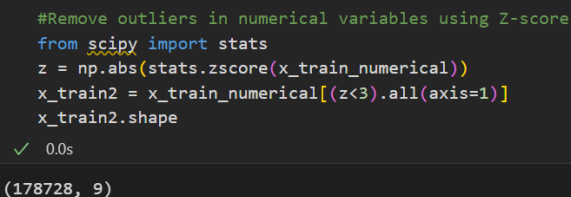
```
Training set size: 178728
Validation set size: 38299
Testing set size: 38300
```

Size of Each Data Set

The dataset is then split into training (70%), validation (15%), and testing (15%) sets. This split ensures the model is trained on a subset of data and evaluate at validation and testing dataset.

Splitting the dataset before preprocessing helps in **avoiding the overfitting** and **data leakage** (Patil, 2024). If model learns patterns from the testing set or the entire dataset, it will result in overfitting. Data leakage can occur when preprocessing steps, such as imputing missing values are applied to the entire dataset before splitting, which will inadvertently introduce information from the testing set into the training process. Therefore, Patil (2024) suggested that splitting dataset before preprocessing will ensure the model is trained solely on patterns and characteristics inherent in the training data, rather than being influenced by specific patterns in the testing set. With that, the model can better generalize to unseen data and provide more reliable performance evaluations.

3.4 Detect and Remove Outlier



```
#Remove outliers in numerical variables using Z-score
from scipy import stats
z = np.abs(stats.zscore(x_train_numerical))
x_train2 = x_train_numerical[(z<3).all(axis=1)]
x_train2.shape

✓ 0.0s
(178728, 9)
```

Check and Remove Outlier

To handle the outlier, we split the 'x_train' variables into numerical and categorical variables. Then, we detect and remove the outlier by using the Z-score method with the threshold of 3. This is because Z-score is sensitive to the spread of data and identifies outlier based on their deviation from mean. However, the result shows that there is no outlier in the dataset.

3.5 Data Transformation and Standardization

Data transformation and standardization are important in preparing data for efficient analysis and modelling. It performs **one-hot encoding** and standardization on dataset to prepare it for

machine learning modelling. The categorical variable is encoded by using `pd.get_dummies()` to convert it into binary variables. Dropping the first category in each categorical variable can avoid multicollinearity issues. This encoding ensures the variable can be used as input for machine learning algorithm. The numerical variable is also **standardized by using `StandardScaler()`** to avoid variable with larger magnitudes from dominating the model. This may standardize the range of features and make the training process more stable. Lastly, all the transformed variable is concatenated for model training.

	Age	Income	LoanAmount	CreditScore	MonthsEmployed	\
34705	-0.034294	-0.815934	0.637581	-0.372791	-0.072250	
91534	1.632722	-1.392594	-0.768511	0.830369	-0.995964	
135742	0.565831	0.140702	-0.573785	-0.139718	1.601982	
38965	1.566041	-0.193464	1.568559	-0.026331	0.995795	
152605	-1.567949	1.032511	-0.639786	0.024063	1.226723	
	NumCreditLines	InterestRate	LoanTerm	DTIRatio	\	
34705	0.445688	-1.406692	-0.707583	1.688273		
91534	1.341111	-0.954764	1.411479	1.688273		
135742	0.445688	-0.475720	-1.413937	1.255511		
38965	-1.345159	0.628490	0.705125	-1.211232		
152605	1.341111	-0.557068	-0.707583	0.043777		
	Education_High	School	...	EmploymentType_Unemployed	\	
34705		False	...	False		
91534		True	...	False		
135742		False	...	True		
38965		True	...	True		
152605		False	...	False		
	MaritalStatus_Married	MaritalStatus_Single	HasMortgage_Yes	\		
34705		True	False	False		
91534		True	False	True		
135742		False	True	False		
...						
38965		False	False	True		
152605		False	False	True		

[5 rows x 24 columns]

Training dataset after encoding and standardization

	Age	Income	LoanAmount	CreditScore	MonthsEmployed	\
249374	1.242256	0.036544	-0.881548	-0.302496	0.731129	
209752	0.974809	0.364488	1.235472	-0.672763	-1.549465	
151150	0.172469	-1.592574	0.746812	-1.538809	0.990943	
215987	-1.499072	1.300999	-1.401759	-0.227188	-0.336997	
162055	-1.164764	0.495491	-0.087641	1.354287	-0.914363	
	NumCreditLines	InterestRate	LoanTerm	DTIRatio	\	
249374	1.348297	-0.275680	-0.707070	1.555027		
209752	-0.443180	0.593941	-1.417833	-0.702950		
151150	-0.443180	-0.841385	0.714456	0.252348		
215987	-0.443180	0.794044	-0.707070	0.686574		
162055	-0.443180	1.510203	1.425219	0.947110		
	Education_High	School	...	EmploymentType_Unemployed	\	
249374		False	...	False		
209752		False	...	False		
151150		False	...	False		
215987		False	...	True		
162055		False	...	False		
	MaritalStatus_Married	MaritalStatus_Single	HasMortgage_Yes	\		
249374		False	True	False		
209752		True	False	True		
151150		False	True	False		
...						
215987		True	False	False		
162055		False	True	True		

[5 rows x 24 columns]

Validation dataset after encoding and standardization

	Age	Income	LoanAmount	CreditScore	MonthsEmployed	\
65844	0.233452	0.619176	-0.553485	-1.304957	-0.738768	
126589	-1.100479	0.482542	-0.878464	0.899993	-0.854218	
229083	-0.033334	0.063220	-1.188118	-1.179319	1.310466	
148419	-0.233424	-0.265123	-1.592375	-0.595101	-1.431467	
25019	1.567382	0.675098	-1.720407	1.660104	1.714541	
	NumCreditLines	InterestRate	LoanTerm	DTIRatio	\	
65844	1.339434	-0.328848	0.699276	1.645913		
126589	0.445520	-0.242718	-0.715492	-1.560416		
229083	0.445520	0.876079	1.406660	-0.390539		
148419	0.445520	-1.295028	0.008108	0.779337		
25019	-0.448376	-0.307693	1.406660	-1.127128		
	Education_High	School	...	EmploymentType_Unemployed	\	
65844		False	...	True		
126589		False	...	True		
229083		False	...	True		
148419		False	...	True		
25019		False	...	False		
	MaritalStatus_Married	MaritalStatus_Single	HasMortgage_Yes	\		
65844		True	False	False		
126589		False	False	False		
229083		False	False	False		
...						
148419		False	False	True		
25019		True	False	False		

[5 rows x 24 columns]

Testing dataset after encoding and standardization

4. Data Mining Model

4.1 Model Training (Without Resampling)

After data preprocessing, we develop seven data mining models to evaluate model and perform model comparison and selection. Below are the seven models with their corresponding performance metrics, confusion matrix and classification report.

Validation Set Metrics: Accuracy: 0.8856889213817594 Precision: 0.6175298804780877 Recall: 0.03493351363533919 F1-score: 0.06612627986348123						Testing Set Metrics: Accuracy: 0.8838642297650131 Precision: 0.610236220472441 Recall: 0.03441385435168739 F1-score: 0.06515342580916351					
Confusion Matrix for validation dataset: [[33766 96] [4282 155]]						Confusion Matrix: [[33697 99] [4349 155]]					
		precision	recall	f1-score	support			precision	recall	f1-score	support
	0	0.89	1.00	0.94	33862		0	0.89	1.00	0.94	33796
	1	0.62	0.03	0.07	4437		1	0.61	0.03	0.07	4504
	accuracy			0.89	38299		accuracy			0.88	38300
	macro avg	0.75	0.52	0.50	38299		macro avg	0.75	0.52	0.50	38300
	weighted avg	0.86	0.89	0.84	38299		weighted avg	0.85	0.88	0.84	38300

Logistic Regression

Validation Set Metrics: Accuracy: 0.8853494869317736 Precision: 0.6197916666666666 Recall: 0.02681992337164751 F1-score: 0.0514149924389717						Testing Set Metrics: Accuracy: 0.8838903394255875 Precision: 0.6446700507614214 Recall: 0.02819715808170515 F1-score: 0.054031057221867684					
Confusion Matrix for validation dataset: [[33789 73] [4318 119]]						Confusion Matrix: [[33726 70] [4377 127]]					
		precision	recall	f1-score	support			precision	recall	f1-score	support
	0	0.89	1.00	0.94	33862		0	0.89	1.00	0.94	33796
	1	0.62	0.03	0.05	4437		1	0.64	0.03	0.05	4504
	accuracy			0.89	38299		accuracy			0.88	38300
	macro avg	0.75	0.51	0.50	38299		macro avg	0.76	0.51	0.50	38300
	weighted avg	0.86	0.89	0.84	38299		weighted avg	0.86	0.88	0.83	38300

Naïve Bayes

Validation Set Metrics: Accuracy: 0.8025535914775843 Precision: 0.1987661461345672 Recall: 0.23236421005183683 F1-score: 0.21425602660016624						Testing Set Metrics: Accuracy: 0.8001044386422976 Precision: 0.1977368622938243 Recall: 0.2289076376554174 F1-score: 0.21218357686766826					
Confusion Matrix for validation dataset: [[29706 4156] [3406 1031]]						Confusion Matrix: [[29613 4183] [3473 1031]]					
		precision	recall	f1-score	support			precision	recall	f1-score	support
	0	0.90	0.88	0.89	33862		0	0.90	0.88	0.89	33796
	1	0.20	0.23	0.21	4437		1	0.20	0.23	0.21	4504
	accuracy			0.80	38299		accuracy			0.80	38300
	macro avg	0.55	0.55	0.55	38299		macro avg	0.55	0.55	0.55	38300
	weighted avg	0.82	0.80	0.81	38299		weighted avg	0.81	0.80	0.81	38300

Decision Tree

Validation Set Metrics:						Testing Set Metrics:					
Accuracy: 0.8861589075432779						Accuracy: 0.8839164490861618					
Precision: 0.6842105263157895						Precision: 0.6559139784946236					
Recall: 0.03222898354744196						Recall: 0.02708703374777975					
F1-score: 0.06155832974601808						F1-score: 0.05202558635394456					
Confusion Matrix for validation dataset:						Confusion Matrix:					
[[33796 66]						[[33732 64]					
[4294 143]]						[4382 122]]					
	precision	recall	f1-score	support			precision	recall	f1-score	support	
0	0.89	1.00	0.94	33862		0	0.89	1.00	0.94	33796	
1	0.68	0.03	0.06	4437		1	0.66	0.03	0.05	4504	
accuracy			0.89	38299		accuracy			0.88	38300	
macro avg	0.79	0.52	0.50	38299		macro avg	0.77	0.51	0.50	38300	
weighted avg	0.86	0.89	0.84	38299		weighted avg	0.86	0.88	0.83	38300	

Random Forest

Validation Set Metrics:						Testing Set Metrics:					
Accuracy: 0.8744092535053134						Accuracy: 0.8742558746736292					
Precision: 0.3184031158714703						Precision: 0.34462151394422313					
Recall: 0.07369844489519946						Recall: 0.07682060390763766					
F1-score: 0.11969253294289897						F1-score: 0.12563543936092955					
Confusion Matrix for validation dataset:						Confusion Matrix:					
[[33162 700]						[[33138 658]					
[4110 327]]						[4158 346]]					
	precision	recall	f1-score	support			precision	recall	f1-score	support	
0	0.89	0.98	0.93	33862		0	0.89	0.98	0.93	33796	
1	0.32	0.07	0.12	4437		1	0.34	0.08	0.13	4504	
accuracy			0.87	38299		accuracy			0.87	38300	
macro avg	0.60	0.53	0.53	38299		macro avg	0.62	0.53	0.53	38300	
weighted avg	0.82	0.87	0.84	38299		weighted avg	0.82	0.87	0.84	38300	

K-Nearest Neighbors

Validation Set Metrics:						Testing Set Metrics:					
Accuracy: 0.8841484111856707						Accuracy: 0.8824281984334204					
Precision: 0.0						Precision: 1.0					
Recall: 0.0						Recall: 0.00022202486678507994					
F1-score: 0.0						F1-score: 0.0004439511653718091					
Confusion Matrix for validation dataset:						Confusion Matrix:					
[[33862 0]						[[33796 0]					
[4437 0]]						[4503 1]]					
	precision	recall	f1-score	support			precision	recall	f1-score	support	
0	0.88	1.00	0.94	33862		0	0.88	1.00	0.94	33796	
1	0.00	0.00	0.00	4437		1	1.00	0.00	0.00	4504	
accuracy			0.88	38299		accuracy			0.88	38300	
macro avg	0.44	0.50	0.47	38299		macro avg	0.94	0.50	0.47	38300	
weighted avg	0.78	0.88	0.83	38299		weighted avg	0.90	0.88	0.83	38300	

Support Vector Machine

Validation Set Metrics:					Testing Set Metrics:				
Accuracy: 0.8861327972009713					Accuracy: 0.8833681462140992				
Precision: 0.5501319261213721					Precision: 0.5257301808066759				
Recall: 0.09398242055442867					Recall: 0.08392539964476022				
F1-score: 0.16053897978825793					F1-score: 0.14474439977024697				
Confusion Matrix for validation dataset:					Confusion Matrix:				
[[33521 341]					[[33455 341]				
[4020 417]]					[4126 378]]				
	precision	recall	f1-score	support		precision	recall	f1-score	support
0	0.89	0.99	0.94	33862	0	0.89	0.99	0.94	33796
1	0.55	0.09	0.16	4437	1	0.53	0.08	0.14	4504
accuracy			0.89	38299	accuracy			0.88	38300
macro avg	0.72	0.54	0.55	38299	macro avg	0.71	0.54	0.54	38300
weighted avg	0.85	0.89	0.85	38299	weighted avg	0.85	0.88	0.84	38300

Neural Network

Based on the results of the seven models, it is evident that there is a significant **class imbalance issue**. Firstly, the **recall scores for class 1 ('default')** which is the minority class **are considerably low**. For class 1, most of the models achieved low recall of 0.03 in validation and testing set, and Support Vector Machine even achieved recall score of 0 in both validation and testing set. Low recall means high false negative, indicating that the models have high chance to miss a large number of actual default applicants. Failing to identify actual default applicants may lead to loss of revenue of the bank authorities and increases the risk of loan defaulted by the applicants. Hence, it is important to have higher recall score in this scenario. Besides, the **confusion matrices for both validation and testing sets of every model appeared to be imbalance**. The training dataset has an uneven distribution of classes, where the number of proportion of instances for class 0 ('non-default') significantly outweighs the class 1 ('default'). The **F1-score for minority class ('default') is observed to be low** across the seven models. Low F1-score means the model has low ability in classifying the applicants correctly to default and non-default categories. Given these reasons, we decided to perform **resampling** on the dataset to address the class imbalance issue and reduce bias towards the majority class ('non-default').

4.2 Data Resampling Technique

To address the issue of imbalanced class, we performed oversampling on the minority class ('default') using **Synthetic Minority Over-sampling Technique (SMOTE)** which generates synthetic samples of minority class (Satpathy, 2020).

```
(array([0, 1], dtype=int64), array([158021, 20707], dtype=int64))
(array([0, 1], dtype=int64), array([158021, 158021], dtype=int64))
```

Size of 'Default' variable in training data set before and after resampling

After resampling, both classes now have 158021 instances, and we can repeat the previous model training process on this resampled dataset.

4.3 Model after Resampling and Hyperparameter Tuning

4.3.1 Logistic Regression

Logistic Regression is widely used for binary classification problems, where the goal is to predict the probability of an observation belongs to either one of two classes for binary classification (Wolff, 2020).

The Logistic Regression is initiated and fitted on training data. We made prediction on validation set and testing set using the trained model to calculate the evaluation metric and confusion matrix. The report for validation set and testing set are shown below.

Validation Set Metrics:							Testing Set Metrics:						
Accuracy: 0.7028643045510327							Accuracy: 0.7066579634464752						
Precision: 0.2049800288943656							Precision: 0.21163567817667722						
Recall: 0.5436105476673428							Recall: 0.5484014209591475						
F1-score: 0.2977042705504814							F1-score: 0.3054095826893354						
Confusion Matrix for validation dataset:							Confusion Matrix:						
[[24507 9355]							[[24595 9201]						
[2025 2412]]							[2034 2470]]						
			precision	recall	f1-score	support				precision	recall	f1-score	support
	0		0.92	0.72	0.81	33862		0		0.92	0.73	0.81	33796
	1		0.20	0.54	0.30	4437		1		0.21	0.55	0.31	4504
	accuracy				0.70	38299		accuracy				0.71	38300
	macro avg		0.56	0.63	0.55	38299		macro avg		0.57	0.64	0.56	38300
	weighted avg		0.84	0.70	0.75	38299		weighted avg		0.84	0.71	0.75	38300

Based on the report, the accuracy for the validation set metrics is 0.70 and for testing set is 0.71 which are moderate. The precision for class 0 ('non-default') is 0.92 for validation and testing set. The precision is 0.20 and 0.21 for class 1 ('default') in validation and testing set. This show that there is a high number of false positives. Besides, the recall for the report is 0.54 and 0.55 for class 1 in validation and testing set while for class 0 is 0.72 and 0.73 respectively. This indicates that the model can capture a good portion of actual positive. However, the F1-score for class 1 in validation and testing set is 0.30 and 0.31 which are considered low while for class 0 is 0.81 for both dataset which is considered good.

```
Best Hyperparameters: {'C': 0.001, 'penalty': 'l2', 'solver': 'liblinear'}
Best Cross-Validation F1 Score: 0.7358528306859768
```

To enhance model simplicity and interpretability, we utilize Grid Search Cross Validation for hyperparameter tuning. The best hyperparameter tuning is 0.001 of regularization strength, L2 penalty and 'liblinear' solver. The corresponding best cross-validation F1-score is 0.74, which is moderate. Then, the Logistic Regression model is trained

again using the best hyperparameters and evaluate the model performance using validation and testing set.

Validation Set Metrics:

Accuracy: 0.695814512128254

Precision: 0.21196390064691317

Recall: 0.5981519044399369

F1-score: 0.313008609505838

Confusion Matrix for validation dataset:

[[23995 9867]

[1783 2654]]

		precision	recall	f1-score	support
	0	0.93	0.71	0.80	33862
	1	0.21	0.60	0.31	4437
	accuracy			0.70	38299
	macro avg	0.57	0.65	0.56	38299
	weighted avg	0.85	0.70	0.75	38299

Testing Set Metrics:

Accuracy: 0.6984856396866841

Precision: 0.21895946377274178

Recall: 0.6092362344582594

F1-score: 0.3221413477342099

Confusion Matrix:

[[24008 9788]

[1760 2744]]

		precision	recall	f1-score	support
	0	0.93	0.71	0.81	33796
	1	0.22	0.61	0.32	4504
	accuracy			0.70	38300
	macro avg	0.58	0.66	0.56	38300
	weighted avg	0.85	0.70	0.75	38300

According to the result after tuning, the **performance of the precision, recall and F1-score increase**. Although, there is a slight decrease in the accuracy in validation and testing set which is 0.70 for both sets. The precision increase to 0.21 and 0.22 for class 1 ('default') in validation and testing set while class 0 ('non-default') precision increase from 0.92 to 0.93 for both sets. This indicate that there is a decrease in number of false positive. The recall for class 1 ('default') increases significantly for both set, which are 0.60 for validation set and 0.61 for testing set. The increase of number show that the model can capture more numbers of actual positive. However, the recall for class 0 ('non-default') in both dataset decreases slightly to 0.71. F1-score for class 1 ('default') is increased to 0.31 and 0.32 for validation and testing set. This indicates a better balance between the precision and recall which is important in classification task especially dealing with imbalance datasets. The improved of performance on both validation and testing set show that the model with tuning is more suitable to make accurate predictions on unseen data. So, we can conclude that the **model after tuning might be more suitable to perform model comparison**.

4.3.2 Naïve Bayes

Naïve Bayes classifier works on Bayes' theorem, where a feature is assumed to be independent from other features. Particularly, Gaussian naive Bayes is used in our model training for classification purposes (Shetty, 2023).

The Naïve Bayes model is instantiated and trained on training data. Prediction is made on validation set and testing set using the trained model to calculate the evaluation metric and confusion matrix. The report for validation set and testing set are shown below.

Validation Set Metrics:							Testing Set Metrics:						
Accuracy: 0.6955272983628815							Accuracy: 0.6981984334203656						
Precision: 0.2126034373010821							Precision: 0.2186328467735503						
Recall: 0.6022086995717827							Recall: 0.6085701598579041						
F1-score: 0.31426051161423113							F1-score: 0.3216947362244						
Confusion Matrix for validation dataset:							Confusion Matrix:						
[[23966 9896]							[[24000 9796]						
[1765 2672]]							[1763 2741]]						
			precision	recall	f1-score	support			precision	recall	f1-score	support	
		0	0.93	0.71	0.80	33862			0.93	0.71	0.81	33796	
		1	0.21	0.60	0.31	4437			0.22	0.61	0.32	4504	
		accuracy			0.70	38299					0.70	38300	
		macro avg	0.57	0.65	0.56	38299			0.58	0.66	0.56	38300	
		weighted avg	0.85	0.70	0.75	38299			0.85	0.70	0.75	38300	

From the evaluation metrics, we can observe that the classifier achieved an accuracy of 0.70 on both the validation set and testing set. This indicates that the model correctly predicted the class label for about 69.5% - 69.8% of the instances in both sets. By looking at the precision of class 1 in both sets, it shows that all the instances predicted as default only about 21% - 22% were actual default loans. Whereas for class 0 ('non-default'), the precision achieves 0.93 for both sets. The recall score for class 0 ('non-default') is 0.71 in both sets while for class 1 ('default') is 0.60 for validation set and 0.61 for testing set. The F1-score for class 1 ('default') is approximately 0.31 on the validation set and 0.32 on the testing set while for class 0 ('non-default') is 0.80 for validation set and 0.81 for testing set. This metric provides a balance between precision and recall. In both sets, the classifier performed better at predicting non-default loans (class 0) than default loans (class 1). Hence, the SMOTE resampling improves the recall and F1-score for class 1 compared to a non-resampled scenario.

4.3.3 Decision Tree

Decision Tree Model which consists of root node, branches, internal and lead nodes, forming a hierarchical tree structure, serves as a non-parametric learning algorithm used for classification and regression tasks (IBM, 2023b).

We instantiated the Decision Tree Classifier and train the model on the training data. We made predictions on the validation set and testing set using the trained model and calculate the evaluation metrics and confusion matrix. The classification reports for validation set and training set are shown below.

Validation Set Metrics:						Testing Set Metrics:					
Accuracy: 0.39951434763309746						Accuracy: 0.4097389033942559					
Precision: 0.13721121144599507						Precision: 0.13908050560229673					
Recall: 0.7910750507099391						Recall: 0.7744227353463587					
F1-score: 0.2338596841894863						F1-score: 0.23581110773079134					
Confusion Matrix for validation dataset:						Confusion Matrix:					
[[11791 22071]						[[12205 21591]					
[927 3510]]						[1016 3488]]					
		precision	recall	f1-score	support			precision	recall	f1-score	support
0	0.93	0.35	0.51	33862		0	0.92	0.36	0.52	33796	
1	0.14	0.79	0.23	4437		1	0.14	0.77	0.24	4504	
accuracy				0.40	38299	accuracy				0.41	38300
macro avg	0.53	0.57	0.37	38299		macro avg	0.53	0.57	0.38	38300	
weighted avg	0.84	0.40	0.47	38299		weighted avg	0.83	0.41	0.49	38300	

Based on the classification report, the accuracy of validation set is 0.40 while for testing set is 0.41 which are quite low. The precision for both classes is low, which is around 0.14 for both datasets, indicating among all the applicants only small proportion were correctly predicted and many non-default applicants were predicted as default applicants. In contrast, the precision for class 0 ('non-default') is high, which is 0.93 for validation set and 0.92 for testing set. Besides, the recall for class 1 ('default') is 0.79 for validation set while 0.77 for testing set, indicating that the model can capture a good portion of actual default applicants. Whereas recall for class 0 ('non-default') is 0.35 for validation set and 0.36 for testing set. The f1-score for class 1 ('default') in validation set and testing set are low which are 0.23 and 0.24 respectively, while the f1-score for class 0 ('non-default') in validation set is 0.51 and testing set is 0.52. This indicates that the model ability to identify applicants who will default on their loans is limited and there might be potential risk to miss true default applicants.

```
Best Hyperparameters: {'criterion': 'entropy', 'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2}
Best Cross-Validation F1 Score: 0.8356757452441291
```

To improve the performance of the model, we perform hyperparameter tuning using Grid Search Cross Validation as Grid Search CV perform exhaustive search across all possible combinations of specified hyperparameters to find the best hyperparameters. The best hyperparameters are where entropy is selected for criterion, maximum depth is set to none, minimum samples leaf is 1 and minimum samples split is 2. The corresponding best cross-validation F1-score is 0.84, which is relatively high. Then, the Decision Tree Model is trained again using the best hyperparameters to evaluate newly trained model's performance using validation set and testing set.

Validation Set Metrics:					Testing Set Metrics:				
Accuracy: 0.31828507271730333					Accuracy: 0.33467362924281985				
Precision: 0.12590623489608507					Precision: 0.1276799659236121				
Recall: 0.8219517692134325					Recall: 0.7986234458259325				
F1-score: 0.21836362003412868					F1-score: 0.22016158648549394				
Confusion Matrix for validation dataset:					Confusion Matrix:				
[[8543 25319]					[[9221 24575]				
[790 3647]]					[907 3597]]				
	precision	recall	f1-score	support		precision	recall	f1-score	support
0	0.92	0.25	0.40	33862	0	0.91	0.27	0.42	33796
1	0.13	0.82	0.22	4437	1	0.13	0.80	0.22	4504
accuracy			0.32	38299	accuracy			0.33	38300
macro avg	0.52	0.54	0.31	38299	macro avg	0.52	0.54	0.32	38300
weighted avg	0.82	0.32	0.38	38299	weighted avg	0.82	0.33	0.40	38300

According to the classification report, the accuracy of both validation and testing set decreases. The precision for class 1 ('default') and class 0 ('non-default') in validation and testing set also decreases by 1% respectively. The recall for class 1 ('default') in validation set and testing set increases and but decreases for class 0 ('non-default'). The f1-score for both classes decreases in both datasets. This indicates that the Decision Tree **model before tuning has better performance**, thus the previous trained model will be used to in comparison later.

4.3.4 Random Forest

A Random Forest Model is a supervised machine learning algorithm. It is based on the combination of multiple decision trees to reach a single result and it is effective in handling the classification task in this assignment (IBM, 2023c).

We instantiated the Random Forest Classifier with a random state of 42 to produce same results across different runs. After training the model on the training data, we made predictions on the validation set and testing set using the trained model and calculate the performance metrics and confusion matrix. The classification reports for validation set and training set are shown below.

Validation Set Metrics:					Testing Set Metrics:				
Accuracy: 0.7072508420585394					Accuracy: 0.7019321148825065				
Precision: 0.20172580787179714					Precision: 0.2059223961878829				
Recall: 0.5163398692810458					Recall: 0.5373001776198935				
F1-score: 0.2901101684183867					F1-score: 0.29773622047244097				
Confusion Matrix for validation dataset:					Confusion Matrix:				
[[24796 9066]					[[24464 9332]				
[2146 2291]]					[2084 2420]]				
	precision	recall	f1-score	support		precision	recall	f1-score	support
0	0.92	0.73	0.82	33862	0	0.92	0.72	0.81	33796
1	0.20	0.52	0.29	4437	1	0.21	0.54	0.30	4504
accuracy			0.71	38299	accuracy			0.70	38300
macro avg	0.56	0.62	0.55	38299	macro avg	0.56	0.63	0.55	38300
weighted avg	0.84	0.71	0.75	38299	weighted avg	0.84	0.70	0.75	38300

Based on the classification report, the accuracy of validation set is 0.71 while for testing set is 0.70 which are moderate. The precision for class 1 ('default') is low, which is around 0.2 for validation set while 0.21 for testing set, indicating a high number of false positives in which many non-default applicants are predicted as default applicants. For class 0 ('non-default'), the precision for validation and testing set is 0.92. Besides, the recall for class 1 ('default') is 0.52 for validation set while 0.54 for testing set, indicating that the model can capture a good portion of actual positive instances where most default applicants are captured correctly. For class 0 ('non-default'), the recall for validation set is 0.73 and testing set is 0.72. However, the f1-score for class 1 ('default') in validation set and testing set are low which are 0.29 and 0.3 respectively. The f1-score for class 0 ('non-default') in validation set is 0.82 while testing set is 0.81.

```
Best Hyperparameters: {'n_estimators': 200, 'min_samples_split': 5, 'min_samples_leaf': 2, 'max_depth': None, 'bootstrap': False}
Best Cross-Validation F1 Score: 0.9301695872977236
```

After that, to explore optimal set of hyperparameters, we performed hyperparameter tuning with Randomized Search Cross Validation as it is scalable to large hyperparameter search spaces with high-dimensional data. The best hyperparameters are 200 n_estimators, 5 minimum number of samples to be split, 2 min_samples_leaf, unlimited depth, and bootstrap is set to false. The corresponding best cross-validation F1-score is 0.93, which is relatively high. Then, we used the best hyperparameters found to instantiate a new Random Forest Classifier and make predictions on the validation set and testing set using the tuned model.

Validation Set Metrics:						Testing Set Metrics:					
Accuracy: 0.7009582495626517						Accuracy: 0.6998955613577024					
Precision: 0.1985736380821447						Precision: 0.2054113283884019					
Recall: 0.5208474194275411						Recall: 0.5410746003552398					
F1-score: 0.28752721617418353						F1-score: 0.2977761485826002					
Confusion Matrix for validation dataset:						Confusion Matrix:					
[[24535 9327]						[[24369 9427]					
[2126 2311]]						[2067 2437]]					
		precision	recall	f1-score	support			precision	recall	f1-score	support
0	0.92	0.72	0.81	33862		0	0.92	0.72	0.81	33796	
1	0.20	0.52	0.29	4437		1	0.21	0.54	0.30	4504	
accuracy				0.70	38299	accuracy				0.70	38300
macro avg	0.56	0.62	0.55	38299		macro avg	0.56	0.63	0.55	38300	
weighted avg	0.84	0.70	0.75	38299		weighted avg	0.84	0.70	0.75	38300	

According to the result after tuning, the overall performance of the metrics remained unchanged except for recall score and f1-score which decreases by 1% for class 0 ('non-default') in validation set. This indicates that the default hyperparameters for the **model before tuning**

contribute to better performance. Thus, we will use the trained Random Forest Model before tuning to perform model comparison.

4.3.5 K-Nearest Neighbors (KNN)

K-Nearest Neighbors (KNN) algorithm is a supervised learning which utilizes proximity to classify or predict the grouping of individual data points (IBM, 2023d).

We instantiated the K-Nearest Neighbors (KNN) Model and trained the model on the training data. Predictions are made on the validation set and testing set using the trained model and calculate the performance metrics and confusion matrix. The classification reports for validation set and training set are shown below.

Validation Set Metrics:

Accuracy: 0.6926029400245437

Precision: 0.18115438108484005

Recall: 0.4696867252648186

F1-score: 0.26146414904962045

Confusion Matrix for validation dataset:

[[24442 9420]

[2353 2084]]

		precision	recall	f1-score	support
	0	0.91	0.72	0.81	33862
	1	0.18	0.47	0.26	4437
	accuracy			0.69	38299
	macro avg	0.55	0.60	0.53	38299
	weighted avg	0.83	0.69	0.74	38299

Testing Set Metrics:

Accuracy: 0.6914099216710182

Precision: 0.17868751647193182

Recall: 0.4515985790408526

F1-score: 0.25605841253855355

Confusion Matrix:

[[24447 9349]

[2470 2034]]

		precision	recall	f1-score	support
	0	0.91	0.72	0.81	33796
	1	0.18	0.45	0.26	4504
	accuracy			0.69	38300
	macro avg	0.54	0.59	0.53	38300
	weighted avg	0.82	0.69	0.74	38300

After SMOTE resampling, the KNN model's performance improved compared to the previous result. The accuracy in both datasets achieved 0.69. However, the precision for class 1 ('default') in both datasets is low, which is 0.18. In contrast, the precision for class 0 ('non-default') is 0.91. In the validation set and testing set, the recall in class 1 ('default') increased significantly to 0.47 and 0.45 respectively, indicating that the model was better at identifying actual default cases. Whereas for class 0 ('non-default'), the recall for both datasets is 0.72. The F1-score also improved after resampling, reaching a value of 0.26 for class 1 ('default') in both datasets and 0.81 for class 0 ('non-default'). This indicates a more balanced performance between precision and recall compared to the initial model. However, the slight drop in accuracy compared with non-resampling may be attributed to the increase in false positives, which indicates that the model is less biased towards predicting non-default cases. Also, by increasing the number of minority class samples, SMOTE can also lead to more false positives, further contributing to the drop in precision.

```
Best Hyperparameters: {'metric': 'euclidean', 'n_neighbors': 4, 'weights': 'uniform'}
Best Cross-Validation F1 score: 0.9060810281750491
```


After that, we performed hyperparameter tuning with Randomized Search Cross validation as well. The best hyperparameters found through the randomized search were Euclidean metric, 4 n_neighbors, and uniform weights, which resulted in a cross-validation F1 score of approximately 0.91. Then, we use the best hyperparameters found to instantiate a new KNN Classifier and make predictions on the validation and testing set.

Validation Set Metrics:					Testing Set Metrics:				
Accuracy: 0.7488707276952401					Accuracy: 0.7483550913838121				
Precision: 0.1890529348217501					Precision: 0.18733252131546893				
Recall: 0.35496957403651114					Recall: 0.34147424511545293				
F1-score: 0.24671052631578946					F1-score: 0.24193802107912538				
Confusion Matrix for validation dataset:					Confusion Matrix:				
[[27106 6756]					[[27124 6672]				
[2862 1575]]					[2966 1538]]				
	precision	recall	f1-score	support		precision	recall	f1-score	support
0	0.90	0.80	0.85	33862	0	0.90	0.80	0.85	33796
1	0.19	0.35	0.25	4437	1	0.19	0.34	0.24	4504
accuracy			0.75	38299	accuracy			0.75	38300
macro avg	0.55	0.58	0.55	38299	macro avg	0.54	0.57	0.55	38300
weighted avg	0.82	0.75	0.78	38299	weighted avg	0.82	0.75	0.78	38300

Based on the result after tuning, accuracy improved in both sets. Even though the recall and f1-score in class 0 ('non-default') and precision in class 1 ('default') are increased for both datasets, its overall performance shows the decrease of recall and F1-score in class 1 ('default') and precision in class 0 ('non-default') for both datasets as well. This indicates that the default hyperparameters for the **model before tuning contribute to better performance**. So, we can conclude that the model before tuning might be more suitable to perform model comparison.

4.3.6 Support Vector Machine

Support Vector Machine (SVM) is a supervised learning algorithm commonly used in machine learning for binary classification and regression tasks by categorizing elements of a dataset into two distinct groups (Tabsharani, 2023).

Support Vector Machine (SVM) Model is instantiated with default kernel – Radial Basis Function (RBF) kernel and trained the model on the training data. Predictions are made on the validation set and testing set using the trained model and calculate the performance metrics and confusion matrix. The classification reports for validation set and training set are shown below.

Validation Set Metrics:

Accuracy: 0.7618475678216141

Precision: 0.23820701989716075

Recall: 0.48027946810908273

F1-score: 0.31846372263319134

Confusion Matrix for validation dataset:

[[27047 6815]

[2306 2131]]

precision

recall

f1-score

support

0

0.92

0.80

0.86

33862

1

0.24

0.48

0.32

4437

accuracy

0.76

38299

macro avg

0.58

0.64

0.59

38299

weighted avg

0.84

0.76

0.79

38299

Testing Set Metrics:

Accuracy: 0.7613838120104439

Precision: 0.24224224224224225

Recall: 0.4835701598579041

F1-score: 0.3227862171174509

Confusion Matrix:

[[26983 6813]

[2326 2178]]

precision

recall

f1-score

support

0

0.92

0.80

0.86

33796

1

0.24

0.48

0.32

4504

accuracy

0.76

38300

macro avg

0.58

0.64

0.59

38300

weighted avg

0.84

0.76

0.79

38300

Based on the classification report, the SVM model's performance improved after using SMOTE resampling technique. The accuracy is 0.76 on both sets. The precision, recall, and F1-score for the minority class improved significantly as before resampling three of them is 0.0. In this case, the precision for class 1 ('default') is 0.24 while for class 0 ('non-default') is 0.92 for both datasets. Similarly, for both datasets, the recall for class 1 ('default') is 0.48 and for class 0 ('non-default') is 0.80. Similar to f1-score, both datasets achieved 0.32 for class 1 ('default') and 0.86 for class 1 ('non-default'). This indicates that the resampled model is better at identifying default loans, which is a critical improvement for a loan default prediction model.

4.3.7 Neural Network

A neural network is a type of machine learning model inspired by the human brain which processes information by simulating the interactions between biological neurons to recognize patterns, assess choices, and make decisions (IBM, 2023a).

Neural Network Model is instantiated using MLPClassifier and trained the model on the training data. Predictions are made on the validation set and testing set using the trained model and calculate the performance metrics and confusion matrix. The classification reports for validation set and training set are shown below.

Validation Set Metrics:

Accuracy: 0.7403587561032925

Precision: 0.22228946041351488

Recall: 0.49673202614379086

F1-score: 0.3071348940914158

Confusion Matrix for validation dataset:

[[26151 7711]

[2233 2204]]

		precision	recall	f1-score	support
	0	0.92	0.77	0.84	33862
	1	0.22	0.50	0.31	4437
accuracy				0.74	38299
macro avg		0.57	0.63	0.57	38299
weighted avg		0.84	0.74	0.78	38299

Testing Set Metrics:

Accuracy: 0.7425848563968669

Precision: 0.22897054357728516

Recall: 0.5022202486678508

F1-score: 0.31453799624556766

Confusion Matrix:

[[26179 7617]

[2242 2262]]

		precision	recall	f1-score	support
	0	0.92	0.77	0.84	33796
	1	0.23	0.50	0.31	4504
accuracy				0.74	38300
macro avg		0.58	0.64	0.58	38300
weighted avg		0.84	0.74	0.78	38300

Based on the classification report, the model's accuracy is 0.74. The precision for both dataset in class 0 ('non-default') is 0.92, whereas for class 1 ('default') is 0.22 in validation set and 0.23 in testing set. For both datasets, the recall for class 0 ('non-default') is 0.77 while for class 1 ('default') is 0.50. Similarly, f1-score for both datasets is 0.84 for 'Non-default' class and 0.31 for 'Default' class.

```
Best Hyperparameters: {'solver': 'adam', 'learning_rate': 'adaptive', 'hidden_layer_sizes': (50, 100, 50), 'alpha': 0.0001, 'activation': 'relu'}
Best Cross-Validation F1 Score: 0.7988424624558816
```

After that, we performed hyperparameter tuning with Randomized Search Cross Validation as well. The best hyperparameters are 'adam' solver, an 'adaptive' learning rate schedule, a three-layer architecture for the hidden layers with 50 neurons in the first and third layers and 100 neurons in the second layer, an L2 penalty of 0.0001 and the 'relu' function. The best cross-validation F1 score is approximately 0.80. This indicates that the model trained with these settings perform well in terms of balancing precision and recall, which crucial for imbalanced classification tasks like loan default prediction.

Validation Set Metrics:

Accuracy: 0.7323428810151701

Precision: 0.20166256157635468

Recall: 0.44286680189317107

F1-score: 0.27713137296382484

Confusion Matrix for validation dataset:

[[26083 7779]

[2472 1965]]

			precision	recall	f1-score	support
	0		0.91	0.77	0.84	33862
	1		0.20	0.44	0.28	4437
	accuracy				0.73	38299
	macro avg		0.56	0.61	0.56	38299
	weighted avg		0.83	0.73	0.77	38299

Testing Set Metrics:

Accuracy: 0.7365013054830287

Precision: 0.20950301517987108

Recall: 0.44738010657193605

F1-score: 0.28537034414388895

Confusion Matrix:

[[26193 7603]

[2489 2015]]


			precision	recall	f1-score	support
	0		0.91	0.78	0.84	33796
	1		0.21	0.45	0.29	4504
	accuracy				0.74	38300
	macro avg		0.56	0.61	0.56	38300
	weighted avg		0.83	0.74	0.77	38300

Based on the result after tuning, the accuracy, precision, recall, and F1-score decrease on both validation and testing set. This indicates that the default hyperparameters for the model before tuning contribute to better performance. Thus, we will use the Neural Network Model before tuning to perform model comparison.

4.4 Comparison and Selection

```
AUC score:
Logistic Regression: 0.7211551489482292
Naive Bayes: 0.7212340625090002
Decision Tree: 0.5677800740289611
Random Forest: 0.686601967357378
K-Nearest Neighbors: 0.6269830579706428
Support Vector Machine: 0.7006866944799084
Neural Networks: 0.7037076247580563
```

AUC Score for 7 Models



ROC Curves for 7 Models

Based on the AUC scores and ROC curves, **Naïve Bayes has the highest AUC** (Area Under the Curve) score among all the seven models which is 0.72. ROC curve above also showed that the Naïve Bayes model performed better compared to other models. This indicates that Naïve Bayes has a **better trade-off between true positive rate and false positive rate**, which will maximize true positive rate while minimize false positive rate. Thus, Naïve Bayes may have better performance in classification of default and non-default loans.

Models	test - F1 score (0)	test - F1 score (1)	AUC Score
Naïve Bayes	0.81	0.32	0.72123
Logistic Regression	0.81	0.32	0.72116
Neural Networks	0.84	0.29	0.70371
Support Vector Machine	0.86	0.32	0.70069
Random Forest	0.77	0.3	0.68660
K-Nearest Neighbors	0.84	0.26	0.62698
Decision Tree	0.42	0.22	0.56778

Comparison for all models using F1-score and AUC score

The table arranges the models based on the descending order of AUC scores, with corresponding F1 scores evaluated on the testing dataset. Based on the table, Decision Tree model has lowest F1 score for both classes and AUC score. Naïve Bayes and Logistic Regression have similar performance, with a highest F1 score for class 1 ('default') and moderately high F1 score for class 0 ('non-default'). The AUC score also performs similar between these two models with Naïve Bayes achieves slightly higher AUC score than Logistic Regression. Even though F1 score for SVM is highest, but the recall is low (0.48) compared with Naïve Bayes and Logistic Regression (0.61). Consider from financial institution

perspective, **minimizing false negatives** (incorrectly classifying actual defaults as non-defaults) **is critical to minimize the financial loss and risk**. As a result, we may prioritize models with higher recall for the default class, even if it comes at the cost of precision. Hence, we **focus on balanced F1 score for both classes and high AUC score** in order to select the best model.

Since Naïve Bayes achieves **relatively high F1 score** for both classes, this indicates a reasonable balance between precision and recall. Besides, Naïve Bayes has the **highest AUC score** among all models, suggesting a good discriminatory power in distinguishing between default and non-default loan applications. Therefore, we can conclude that **Naïve Bayes model is the best model**.

5. Interaction stage

After selecting the best model, we store the Naïve Bayes model using the Joblib library. This step is to ensure that we can retrieve and use the model without retrain it every time. After that, we load the model into the memory to store it for prediction on new data. Then, the new data which is the new applicants' information is evaluated by using the training model that we have loaded in memory. The new data undergoes similar data preprocessing before interacting with the trained model. This stage involves the separating features from the target variable ('Default' variable) and identifying the categorical and numerical variables. The encoding of categorical variables and standardization of numerical variables are carried out so that data representation is consistent with the trained model in order to get accurate prediction. After that, we combine all the dataset together for model to make prediction. Using the loaded model, we predict the result for the new dataset.

LoanID	Default
A01	0
A02	0
A03	1
A04	0
A05	0
A06	1
A07	1
A08	0
A09	0
A10	0
B01	0
B02	0
B03	0
B04	1
B05	0
B06	0
B07	1
B08	0
B09	1
B10	0

Prediction on Loan Default Status of 20 New Applicants

The prediction for the loan default status is shown as figure above. Out of 20 applicants, there are **14 applicants predicted as non-defaulters (class 0)** and **6 applicants are predicted to be defaulters (class 1)**. This means **14 of the non-defaulters will get loan approval** from bank while the **other 6 applicants will get rejected on their loan requests**. By identifying potential defaulters, our model can effectively mitigate financial risk for bank authorities. This proactive approach allows banks to allocate resources more efficiently and make informed decisions regarding loan approvals. Additionally, for loan applicants, the model increases the chances for non-defaulters to get approved for their loan requests. By accurately assessing risk, the model ensures that deserving applicants are more likely to receive loan approvals. Thus, utilizing this machine learning model on predicting the loan default status is useful to mitigate financial risk and facilitating the loan processes.

6. References

IBM (2023a). *What Are Neural Networks?* / IBM. [online] [www.ibm.com](https://www.ibm.com/topics/neural-networks). Available at: <https://www.ibm.com/topics/neural-networks>.

IBM (2023b). *What is a Decision Tree* / IBM. [online] [www.ibm.com](https://www.ibm.com/topics/decision-trees). Available at: <https://www.ibm.com/topics/decision-trees>.

IBM (2023c). *What is Random Forest?* / IBM. [online] [www.ibm.com](https://www.ibm.com/topics/random-forest). Available at: <https://www.ibm.com/topics/random-forest>.

IBM (2023d). *What is the k-nearest neighbors algorithm?* / IBM. [online] [www.ibm.com](https://www.ibm.com/topics/knn). Available at: <https://www.ibm.com/topics/knn>.

Patil, M. (2024). *The Roles of Data Sets in Machine Learning Projects: A Guide to Data Splitting*. [online] Medium. Available at: <https://medium.com/@madhuri15/the-roles-of-data-sets-in-machine-learning-projects-a-guide-to-data-splitting-454beb468a49>.

Satpathy, S. (2020). *SMOTE / Overcoming Class Imbalance Problem Using SMOTE*. [online] Analytics Vidhya. Available at: <https://www.analyticsvidhya.com/blog/2020/10/overcoming-class-imbalance-using-smote-techniques/#:~:text=SMOTE%20is%20an%20oversampling%20technique>.

Shetty, B. (2023). *An in-depth guide to supervised machine learning classification*. [online] Built In. Available at: <https://builtin.com/data-science/supervised-machine-learning-classification>.

Tabsharani, F. (2023). *What is a support vector machine? / Definition from WhatIs*. [online] WhatIs.com. Available at: [https://www.techtarget.com/whatis/definition/support-vector-machine-SVM#:~:text=A%20support%20vector%20machine%20\(SVM\)%20is%20a%20type%20of%20supervised](https://www.techtarget.com/whatis/definition/support-vector-machine-SVM#:~:text=A%20support%20vector%20machine%20(SVM)%20is%20a%20type%20of%20supervised).

Wolff, R. (2020). *Classification Algorithms in Machine Learning: How They Work*. [online] MonkeyLearn Blog. Available at: <https://monkeylearn.com/blog/classification-algorithms/>.